

ReSearch Flow — AI-Powered Academic Research Assistant

Final Year Project Documentation & Evaluation Report

Field	Details
Project Title	ReSearch Flow: An AI-Powered Academic Research Assistant
University	University of the Punjab, Lahore
Department	Data Science, FCIT
Supervisor	Dr. Khurram Shahzad
Academic Year	2024–2025

Team Members

Roll No.	Name	Role
BSDSF22M054	Shahzeb Ali	Backend Development & System Architecture
BSDSF22M048	Areesha Rizwan	Database Administration & Schema Design
BSDSF22M020	Haroon Mahmood	Web Deployment & DevOps
BSDSF22M028	Maryam Abid	Frontend Development & UI/UX

1. Executive Summary

ReSearch Flow is a full-stack, AI-powered academic research platform built as a Final Year Project for the BS Data Science programme at the University of the Punjab. The system solves a core problem faced by every researcher: the fragmented, time-consuming process of discovering, reading, and synthesising scholarly literature.

The platform aggregates research papers from **five distinct academic sources** (OpenAlex, arXiv, PubMed Central, CORE, and Semantic Scholar), applies intelligent **Title-Based Fuzzy Deduplication** across results, visualises inter-paper relationships using a real-time **Citation Mesh** powered by graph-theoretic algorithms (PageRank + In-Degree Centrality), and leverages **Large Language Models** via OpenRouter to produce structured, section-wise summaries of full-length PDF documents.

Key Achievements

- Multi-source parallel API integration with unified schema normalisation
- Custom PageRank implementation for citation prestige scoring
- Resilient AI summarisation with fallback to abstract-based summarisation for paywalled content
- Interactive, animated Cytoscape.js citation network with bridge-node detection
- Full user authentication (Email + Google OAuth) with per-user saved paper libraries
- Real-time search history and trending topic suggestions powered by OpenAlex




Figure 1: The ReSearch Flow landing page featuring the intelligent search bar with semantic query detection, live trending topics, and key statistics.

2. Problem Statement & Motivation

Academic literature review is one of the most labour-intensive phases of any research project. Researchers typically: 1. Search across multiple disconnected databases (Google Scholar, PubMed, arXiv, etc.) 2. Manually remove duplicate results that appear across sources 3. Read full papers to extract key ideas — often 20–40 pages each 4. Manually trace citations to find foundational or related work 5. Struggle to identify which papers are truly *influential*


vs. merely frequently cited

ReSearch Flow addresses all five pain points in a single, unified interface. The platform's value proposition is summarised below:

Pain Point	ReSearch Flow Solution
Multiple databases to search	Single search bar queries 5 APIs in parallel
Duplicate results	Title-based fuzzy deduplication via Pandas
Reading full papers	AI-generated structured summaries (Problem → Conclusion)
Manual citation tracing	Automated Citation Mesh with depth control
Identifying influential work	PageRank + In-Degree Centrality scores on every node

3. System Architecture

ReSearch Flow adopts a **3-Tier N-Layer Architecture** based on MVC principles, ensuring clean separation of concerns and horizontal scalability.



3.1 Request Lifecycle

1. User enters a topic on the frontend
2. React calls GET /api/v1/search/fetch?topic=<query>

- 3. FastAPI fires parallel requests to OpenAlex, arXiv, PMC, CORE, Semantic Scholar
- 4. Each source's response is normalised to the unified Paper schema (Pydantic)
- 5. Results are merged, deduplicated, and returned immediately to the frontend
- 6. A background task (BackgroundTasks) asynchronously upserts results into Supabase
- 7. Frontend renders the paper list; citation graph and summaries are fetched on demand

4. Technology Stack

4.1 Frontend

Technology	Version	Purpose
React	19	Component-based UI framework
Vite	Latest	Fast dev server & build tool
Tailwind CSS	4	Utility-first styling
React Router	v7	Client-side routing
Cytoscape.js	Latest	Citation network graph rendering
Framer Motion	Latest	Smooth UI animations
Lucide React	Latest	Icon library
Three.js	Latest	Animated 3D background
jsPDF	Latest	Client-side PDF export
xlsx	Latest	Excel export
Supabase JS	Latest	Auth & database client

4.2 Backend

Technology	Version	Purpose
Python	3.12	Core runtime

Technology	Version	Purpose
FastAPI	Latest	Async REST API framework
Uvicorn	Latest	ASGI server
Pydantic	v2	Data validation & schemas
Pandas	Latest	Data cleaning & deduplication
Requests / aiohttp	Latest	HTTP clients for external APIs
PyMuPDF (fitz)	Latest	PDF text extraction
pdfplumber	Latest	Alternative PDF parser
pydantic-settings	Latest	Environment config management

4.3 Database & Auth

Technology	Purpose
Supabase (PostgreSQL)	Primary relational database
Supabase Auth	Email/password + Google OAuth
Row Level Security (RLS)	Per-user data isolation

4.4 External API Integrations

API	Purpose	Auth Required
OpenAlex	Primary paper search (free, no key)	No
arXiv API	Physics, CS, Math papers	No
EuropePMC	Biomedical papers (PubMed Central)	No
CORE API	Open access aggregator	API Key
Semantic Scholar	Citation graph data	No

API	Purpose	Auth Required
OpenRouter	LLM gateway (GPT-4o-mini)	API Key
SerpAPI	Google Scholar scraping	API Key

5. Project Structure

FYP-ReSearch-Flow/ | ├── backend/ | ├── src/app/ | | ├── main.py # FastAPI entry point, CORS, lifespan | | ├── api/v1/ | | | ├── api_router.py # Central route registry | | | └── routes/ | | | ├── search.py # Primary search endpoint | | | ├── papers.py # CRUD for saved papers | | | ├── citation.py # Citation network builder | | | ├── arxiv.py # arXiv: search/summarize/ask | | | ├── core.py # CORE API integration | | | ├── pmc.py # PubMed Central integration | | | ├── semantic_scholar.py | | | └── google_scholar.py | | ├── core/ | | | ├── config.py # Pydantic settings (env vars) | | | ├── logging_config.py | | | ├── crud/ | | | └── papers_crud.py # Supabase upsert logic | | ├── db/ | | | └── client_supabase.py # Supabase client singleton | | ├── models/ | | | └── pydantic_schemas.py # Paper data model | | ├── services/ | | | ├── openalex_service.py | | | ├── citation_network_service.py # PageRank + graph builder | | | ├── llm_service.py # OpenRouter/OpenAI calls | | | ├── summarization_service.py | | | ├── chat_service.py # Ask-about-paper Q&A | | | ├── pdf_service.py | | | ├── pdf_extractor.py | | | ├── arxiv_service.py | | | ├── core_service.py | | | ├── pmc_service.py | | | └── semantic_scholar_service.py | | └── utils/ | | ├── dedupe.py # Pandas-based deduplication | | ├── retries.py # Exponential backoff | └── requirements.txt | ├── frontend/src/ | ├── pages/ | | ├── Home.jsx # Landing page & search | | ├── Results.jsx # 3-column results layout | | ├── SavedSearches.jsx # Bookmarked papers | | └── AuthPage.jsx # Login/Signup | ├── components/ | | ├── SearchBar.jsx # Smart search w/ suggestions | | ├── CitationMesh.jsx # Cytoscape.js graph | | ├── Navbar.jsx | | └── ArticleNotesModal.jsx | ├── contexts/AuthContext.jsx # Global auth state | ├── hooks/useSearchHistory.js # Search history + trending | └── services/api.js # All API calls | ├── database/ | ├── saved_articles_schema.sql | ├── search_history_schema.sql | ├── docs/ # All documentation | ├── scripts/ # Utility scripts | ├── run-servers.ps1 # One-click startup | └── .env.example # Environment template

6. Core Modules — Deep Dive

6.1 Search & Multi-Source Aggregation

The primary search endpoint (`/api/v1/search/fetch`) orchestrates parallel queries to all configured data sources. Each source has a dedicated service module that:

1. Constructs the source-specific API request
2. Parses the response into the unified Paper Pydantic schema
3. Handles rate-limiting, timeouts, and partial failures gracefully

Unified Paper Schema (pydantic_schemas.py): `python class Paper(BaseModel): paperId: Optional[str] title: str abstract: Optional[str] authors: List[str] = [] url: Optional[str] year: Optional[int] venue: Optional[str] citationCount: Optional[int] referenceCount: Optional[int] referencedWorks: List[str] = [] isOpenAccess: Optional[bool] openAccessPdf: Optional[str] externalIds: Dict = {} fieldsOfStudy: List[str] = [] institutions: List[str] = [] source: Optional[str] topic: Optional[str] content: Optional[str]`




Figure 2: The Results page showing aggregated papers from multiple sources. The three-column layout displays the paper list (left), detailed paper view (centre), and statistics/filters (right).

6.2 Data Deduplication (dedupe.py)

After collecting results from multiple APIs, the same paper may appear under slightly different titles (e.g., from arXiv and OpenAlex simultaneously). The `clean_and_deduplicate()` function:

```
python def clean_and_deduplicate(df, topic): # Normalise string fields df['title'] = df['title'].fillna('N/A').astype(str).str.strip() df['abstract'] = df['abstract'].fillna('N/A').astype(str).str.strip() # Remove exact-title duplicates df = df.drop_duplicates(subset=['title'], keep='first') return df
```

Key steps: - **Schema enforcement:** Ensures all expected columns exist, fills missing values - **Type coercion:** Numeric fields cast with `pd.to_numeric(..., errors='coerce')` - **Deduplication:** Title-based, keeping the entry with the richest source data - **List fields preserved:** Authors, institutions, and referenced works remain as Python lists (stored as JSONB in Supabase)

6.3 Intelligent Citation Network

The `citation_network_service.py` module is the most algorithmically rich component of the system. It builds a **bidirectional citation graph** and computes three graph metrics:

Graph Construction

The service auto-detects the paper source (Semantic Scholar IDs are 40-char hex strings; OpenAlex IDs start with w). It then: 1. Adds all searched papers as **root nodes** 2. Concurrently fetches references (outgoing edges) and citing papers (incoming edges) using `ThreadPoolExecutor(max_workers=5)` 3. Identifies **bridge nodes** — non-root papers connected to 2+ different root papers (visually highlighted in orange)

Graph Metrics

1. Citation Velocity (recency-weighted impact): `python age = max(1, current_year - year)`
`node["citationVelocity"] = round(citations / age, 2)`

2. In-Degree Centrality (direct influence score): `python node["influenceScore"] = round(in_degrees[nid] / max_in_degree, 4)`

3. PageRank (global prestige, 10 iterations, damping factor $d=0.85$): `python d = 0.85 for _ in range(10): for node in nodes: pr_sum = sum( node_map[src]["pageRankScore"] / out_degrees[src] for src in incoming_edges[node["id"]] if out_degrees[src] > 0 ) new_pr[node["id"]] = ((1.0 - d) / num_nodes) + d * pr_sum`

These scores are normalised to $[0, 1]$ and sent to the frontend where Cytoscape.js maps them to **node size** and **colour intensity**.




Figure 3: The Citation Mesh showing an interactive graph. Node size encodes PageRank (prestige), node colour encodes In-Degree Centrality (influence), and orange nodes are bridge papers connecting multiple clusters.

6.4 AI Summarisation & Document Q&A

The LLM pipeline (`llm_service.py`) accepts full paper text and produces structured summaries in four sections: Problem Statement, Methodology, Key Findings, and Conclusion.

Pipeline: 1. Strip references/bibliography section (regex-based) to save tokens 2. Truncate to 100,000 characters (fits within GPT-4o-mini's 128K context) 3. Call OpenRouter API with `response_format: json_object` to enforce structured output 4. Parse JSON response; fall back to regex-based text extraction if JSON fails 5. If PDF extraction itself fails (paywalled domain), fall back to abstract-based summarisation

Prompt Engineering: You are an expert academic researcher. Analyze the following research paper and provide a detailed, structured summary in EXACTLY this JSON format: { "problem_statement": "...", "methodology": "...", "key_findings": "...", "conclusion": "..." }

Document Q&A (`chat_service.py`): The same PDF text is passed as system context, and the user's question is appended as the user message. Conversation history is maintained on the frontend and passed back on each turn for multi-turn dialogue.

6.5 Frontend Architecture

The React frontend is structured around three main pages:

Home Page (Home.jsx)

- Clean hero layout with the smart SearchBar component
- Statistics row: 10M+ papers, 500K+ researchers, 99.9% uptime
- Three feature cards: Smart Search, Auto Summarisation, Citation Networks

Results Page (Results.jsx)

The most complex component — a three-column layout: - **Left panel:** Paginated list of paper cards with source badges, year, citation count - **Centre panel:** Selected paper's full detail view — abstract, authors, institutions, AI summary, and Ask-Paper chat - **Right panel:** Real-time statistics, year distribution chart, field-of-study breakdown, and filter controls

Search Bar (SearchBar.jsx)

A sophisticated search component with: - **Semantic query detection:** Regex patterns detect natural language questions and activate smart keyword extraction - **Tabbed dropdown:** "Suggestions" (search history + topic autocomplete) and "Trending" (live from OpenAlex) - **Keyboard navigation:** Arrow keys, Tab to switch tabs, Enter to search, Escape to close - **Debounced API calls:** 300ms for suggestions, 500ms for trending topics

6.6 Database Design

research_papers **Table**

The main paper store. Uses JSONB for flexible array/object fields:

Column	Type	Description
paperId	TEXT (PK)	Unique paper identifier
title	TEXT	Paper title
abstract	TEXT	Paper abstract
authors	JSONB	Array of author names
url	TEXT	Link to paper
year	INTEGER	Publication year

Column	Type	Description
venue	TEXT	Journal/conference name
citationCount	INTEGER	Number of citations
referencedWorks	JSONB	List of cited paper IDs
isOpenAccess	BOOLEAN	Open access status
openAccessPdf	TEXT	Direct PDF URL
fieldsOfStudy	JSONB	Research domains
institutions	JSONB	Author institutions
source	TEXT	API source (openalex, arxiv, etc.)
topic	TEXT	Search topic

saved_articles Table

Per-user bookmarks with Row Level Security:

```
sql CREATE TABLE saved_articles ( id UUID PRIMARY KEY DEFAULT gen_random_uuid(), user_id UUID NOT NULL REFERENCES auth.users(id) ON DELETE CASCADE, paper_id TEXT NOT NULL, title TEXT NOT NULL, notes TEXT, created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(), updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(), UNIQUE(user_id, paper_id) ); ALTER TABLE saved_articles ENABLE ROW LEVEL SECURITY; -- RLS: Users can only access their own rows CREATE POLICY "own_articles" ON saved_articles USING (auth.uid() = user_id);
```

An auto-trigger updates updated_at on any modification.

7. API Endpoints Reference

Method	Endpoint	Description
GET	/health	Backend health check
GET	/api/v1/search/fetch	Search papers (topic, limit, sources)
GET	/api/v1/papers	Retrieve stored papers by topic

Method	Endpoint	Description
POST	/api/v1/citation/network	Build citation graph for paper list
GET	/api/v1/arxiv/search	Search arXiv
POST	/api/v1/arxiv/summarize	AI summarise an arXiv paper
POST	/api/v1/arxiv/ask	Q&A on an arXiv paper
GET	/api/v1/pmc/search	Search PubMed Central
POST	/api/v1/pmc/summarize	AI summarise a PMC paper
POST	/api/v1/pmc/ask	Q&A on a PMC paper
GET	/api/v1/core/search	Search CORE API
POST	/api/v1/semantic-scholar/search	Search Semantic Scholar

8. Authentication & Security


Authentication is handled entirely by **Supabase Auth**, supporting: - **Email/Password** signup and login with email verification - **Google OAuth** via Supabase's OAuth provider integration

The frontend stores session tokens in Supabase's built-in session management. The `AuthContext.jsx` React context exposes `user`, `session`, `signIn`, `signOut`, and `signUp` to all components.

Security measures: - Row Level Security (RLS) ensures users can only read/write their own saved articles - API keys are stored in `.env` files (not committed to Git) - CORS is restricted to the configured frontend origin - All backend endpoints perform schema validation via Pydantic

9. Data Flow Diagrams

9.1 Search Flow



9.2 Citation Network Flow

User clicks "Citation Mesh" | ▼ POST /api/v1/citation/network citation_network_service.py | |→ ThreadPoolExecutor (5 workers) | |— fetch references (outgoing edges) | |— fetch citations (incoming edges) | |→ Compute PageRank (10 iterations) |→ Compute In-Degree Centrality |→ Compute Citation Velocity |→ Identify Bridge Nodes | ▼ Return {nodes, edges, stats} Cytoscape.js (CitationMesh.jsx) |— Render interactive graph

9.3 AI Summarisation Flow

User clicks "Summarise" | ▼ POST /api/v1/arxiv/summarize (or pmc/core) pdf_extractor.py |— Try: Download & extract PDF text (PyMuPDF) |— Fallback: Use title + abstract if domain is blocked | ▼ llm_service.py |— Strip references section |— Truncate to 100K chars |— Call OpenRouter API (GPT-4o-mini) | ▼ JSON response Parse structured summary {problem_statement, methodology, key_findings, conclusion} | ▼ Return to Frontend → Display in Detail Panel

10. Setup & Deployment Guide

Prerequisites

- Python 3.12+
- Node.js 18+
- A Supabase project (free tier works)
- OpenRouter API key (for AI summarisation)

Quick Start

Step 1: Clone & configure bash git clone https://github.com/Shahzeb-Ali-Web-Developer/FYP-ReSearch-Flow.git cd FYP-ReSearch-Flow cp .env.example backend/.env # Fill in your API keys

Step 2: Backend bash cd backend python -m venv venv .\venv\Scripts\Activate.ps1 pip install -r requirements.txt

Step 3: Frontend bash cd frontend npm install

Step 4: Launch both servers powershell .\run-servers.ps1

Verify: - Frontend: http://localhost:5173 - Backend API: http://localhost:8000/health - Interactive API Docs: http://localhost:8000/docs

Environment Variables

```
backend/.env: SUPABASE_URL=https://xxxx.supabase.co SUPABASE_KEY=your_service_role_key
OPENROUTER_API_KEY=your_openrouter_key CORE_API_KEY=your_core_key # Optional
SERPAPI_API_KEY=your_serpapi_key # Optional

frontend/.env: VITE_API_URL=http://127.0.0.1:8000/api/v1 VITE_SUPABASE_URL=https://xxxx.supabase.co
VITE_SUPABASE_ANON_KEY=your_anon_key
```

11. Features Summary

Feature	Status	Module
Multi-source paper search	✔ Complete	search.py, all *_service.py
Real-time fuzzy deduplication	✔ Complete	dedupe.py
Paper details panel	✔ Complete	Results.jsx
Advanced filters (year, type, OA)	✔ Complete	Results.jsx
Statistics dashboard	✔ Complete	Results.jsx
AI Summarisation (PDF)	✔ Complete	llm_service.py
Paywall fallback summarisation	✔ Complete	summarization_service.py
Ask About This Paper (Q&A)	✔ Complete	chat_service.py
Citation Mesh (PageRank)	✔ Complete	citation_network_service.py
Bridge node detection	✔ Complete	citation_network_service.py
Citation Velocity metric	✔ Complete	citation_network_service.py
User authentication (Email + OAuth)	✔ Complete	AuthContext.jsx
Save & bookmark papers	✔ Complete	papers_crud.py
Per-user notes on papers	✔ Complete	ArticleNotesModal.jsx
Export to PDF / Excel	✔ Complete	Results.jsx
Search history & trending	✔ Complete	useSearchHistory.js

Feature	Status	Module
Semantic NLP query detection	✔ Complete	SearchBar.jsx

12. Conclusion & Future Work

ReSearch Flow successfully demonstrates a production-grade, full-stack research platform built entirely by a four-person undergraduate team. The project goes well beyond basic API integration by implementing original algorithmic work:

- A **custom PageRank** implementation tuned for academic citation graphs
- A **multi-threaded concurrent graph builder** using Python's `ThreadPoolExecutor`
- A **resilient LLM pipeline** with automatic fallback strategies
- A **bridge-node detection algorithm** for identifying cross-cluster connectors

Future Enhancements

1. **Semantic search:** Replace keyword search with embedding-based vector similarity
2. **Recommendation engine:** Suggest papers based on user reading history
3. **Collaborative libraries:** Shared paper collections for research groups
4. **Docker containerisation:** One-command deployment via Docker Compose
5. **CI/CD pipeline:** GitHub Actions for automated testing and deployment
6. **Mobile PWA:** Progressive Web App for mobile access